

Reading orientation. Verification companion to "The η Corridor in 3D Wilson-Fisher Critical Phenomena: Empirical Observation and Structural Architecture" (D.B., 2026), OSF DOI: 10.17605/OSF.IO/PD73B (shared with the main paper – single OSF project). To be read alongside the main paper; the main paper is self-contained without this companion, but this companion presumes familiarity with the main paper's notation and theorem statements.

NOTE · The η Corridor – Verification Companion

Verification Companion to "The η Corridor in 3D Wilson-Fisher Critical Phenomena: Empirical Observation and Structural Architecture"

Companion document – high-precision numerical verification of algebraic identities and empirical fits in the main paper

Published 2026-06-02 · DOI: 10.17605/OSF.IO/PD73B · OSF: osf.io/pd73b · License: CC BY 4.0 (DOI and OSF shared with the main paper – single OSF project) · Reprinted 16 September 2026

■ READING UPDATE – 2026-09-16. No value in this paper is changed by this notice.
How the corpus is now read: One coordinate up – <https://osf.io/ujwc7>
DOI 10.17605/OSF.IO/UJWC7

Reprint note, 2026-09-11: every table in this companion re-computed at 40–160 digits (g on the test grid, $d(\ln g)/du$, the Mellin values $M[g](s)$ for $s = 1.5$, $g(5/4) = 4/15$, $K(n)$ for $n = 2..10$ and $K'(2)$, the five hyperbolic values at $\ln(2)/2$, $1/\ln(2) - \sqrt{2}$). Five printed strings carried trailing digits beyond what their own computation had earned, and are corrected in place: 2G (§2; its 100th–101st decimals were dps-100 noise), $M[g](2)$ and $M[g](5)$ (§9.9 table, last printed digit of a dps-50 quadrature), and the displayed $\sinh(\ln 2/2)$ and $\cosh(\ln 2/2)$ (§12.2 output, wrong from the 46th decimal – the run that produced the display was not the dps-100 run the caption names; the identities and the residuals are unaffected). Every identity, every residual and every value to the digits actually earned stands. Caught by TOOLBOX/printed_digits.py at --min-dp=7 on the reprint; the lesson (print only the digits the depth earned) is the substrate rule of 2026-09-11.

Created: 2026-05-23 (Saturday very late night) – alongside the main paper draft Companion to: Eta_Observation_Paper_Draft.md Author: D.B. (with Cowork-Opus as drafting collaborator)
Verification tooling: Python 3, mpmath library version 1.3.0 Default precision: dps = 100 (decimal places) unless otherwise stated; cross-architecture verification independently performed at dps = 1050 (see §3).

1. Purpose and scope

This companion document re-verifies, at high arbitrary-precision arithmetic, every algebraic identity and empirical fit stated as a theorem, proposition, or model in the main paper. Each verification is paired with a paper-section cross-reference. The companion is a re-derivation artifact: any reader with a Python environment and the mpmath library can reproduce these computations and confirm that the paper's claims hold at the stated precisions.

The companion is supplementary. The paper's mathematical argument is self-contained without it. The companion exists so that:

- A reviewer can verify the algebraic identities of Parts II §§9–13 to whatever precision they require.
- A future researcher building on the framework can confirm the computational substrate without re-deriving every identity from scratch.
- The framework's verification history (cross-architecture, multi-precision) is preserved as documented evidence rather than as a claim.

What this companion does not contain:

- The derivations themselves (those are in the paper).
- The empirical critical-exponent values from the literature (cited in Part I §5 of the paper, with citations to Chester et al. 2020, Kos et al. 2016, Hasenbusch 2023, Mudrov–Varnashev 2001, Calabrese–Martin-Mayor–Pelissetto–Vicari 2003, Hasenbusch et al. 2007, and the Clisby–Dünweg 2016 (v) + Clisby 2017 (γ) pair from which the SAW η is derived via Fisher's relation $\eta = 2 - \gamma/\nu$).
- The framework's broader architectural locks (those live in the framework atlas, a separate artifact mentioned briefly in §5 below).

2. Setup

The following preamble is assumed at the start of all verification blocks below.

```
from mpmath import mp, mpf, sqrt, log, sinh, cosh, tanh, sech, asinh, acosh, atanh, exp, quad
import mpmath
mp.dps = 100

# Framework constants
K_AUD = sqrt(2) * log(2)
```

```

G      = 1 - K_AUD
two_G  = 2 * G
phi    = (1 + sqrt(5)) / 2

```

Numerical values at dps = 100:

```

K_AUD = sqrt(2)*ln(2) = 0.980258143468547191713901723635233381291460699099054721042246
24706529
10985142058941430135340632871525
G      = 1 - K_AUD      = 0.019741856531452808286098276364766618708539300900945278957753
75293470
890148579410585698646593671284746
2G     = 0.039483713062905616572196552729533237417078601801890557915507
50586941
780297158821171397293187342569491
phi    = (1+sqrt(5))/2 = 1.618033988749894848204586834365638117720309179805762862135448
62270526
0462818902449707207204189391137

```

3. Verifications

Note on precision (dps) per verification block. The verifications below use different dps values depending on the method used in each block, not on any difference in the framework's verification standard:

- Algebraic identity checks (Verifications 9.1, 9.2, 9.4, 9.7, 9.10, 10.1, 10.2, 11.1, 12.1, 12.2, 13.1): closed-form-to-closed-form or symbolic-trace comparisons. These are computationally cheap at any dps. The default precision is dps = 100, with residuals at the working-precision floor ($\sim 10^{-102}$ or exact zero). Theorem 9.7 (closed-form Mellin spectrum) is additionally re-verified at dps = 1000 (see Re-verification 9.7-R), matching the framework's higher-precision standard for the paper's central published theorem.
- Quadrature-based checks (Verifications 9.5 dome integral identity, 9.6 Mellin decomposition with J(s) residual): involve numerical integration via `mpmath.quad` over $[0, \infty)$. Quadrature at high dps converges slowly – adaptive sampling needs exponentially more points to reach the requested precision. The dps values chosen here (dps = 80 for Verification 9.5, dps = 50 for Verification 9.6) are the practical compute floor for those quadratures on commodity hardware; the residuals are at the working-precision floor of the quad operation at that dps. Higher precision is achievable but takes minutes-to-hours rather than seconds; the residuals at the chosen dps already confirm the identities to the precision floor of the numerical integrator.

The cross-architecture verification history (§4) records independent re-verifications at higher precision (dps up to 1500) for the algebraic identities; those higher-precision runs confirm the same identities and are documented per architecture and date.

Verification 9.1 – The sech^2 / doubling-rate factorization

Paper reference: Part II §9.2, Theorem 9.1.

Claim. $g(x) = \tanh(x) / \cosh(2x) = (1/2) \cdot \text{sech}^2(x) \cdot \tanh(2x)$ for all $x \geq 0$.

Code.

```

def g_direct(x):
    return tanh(x) / cosh(2*x)

def g_factored(x):
    return mpf("0.5") * sech(x)**2 * tanh(2*x)

test_x = [mpf("0.1"), mpf("0.3"), mpf("0.5"), mpf("0.7"),
          mpf("1.0"), mpf("1.5"), mpf("2.0"), mpf("3.0")]
for x in test_x:
    a, b = g_direct(x), g_factored(x)
    print(f"x = {x}: residual = {abs(a-b)}")

```

Output (dps = 100).

x	g_direct	g_factored	residual
0.10	0.0977073255999481836198...	0.0977073255999481836198...	8.930e-103
0.30	0.2457369545464443800436...	0.2457369545464443800436...	1.786e-102
0.50	0.2994769986957551296171...	0.2994769986957551296171...	7.144e-102
0.70	0.2809838710850990112747...	0.2809838710850990112747...	3.572e-102
1.00	0.2024334241200519958269...	0.2024334241200519958269...	1.786e-102
1.50	0.0899065000418640130895...	0.0899065000418640130895...	0.000e+00
2.00	0.0353017196632501598458...	0.0353017196632501598458...	0.000e+00
3.00	0.0049329579640651192324...	0.0049329579640651192324...	5.581e-104

Maximum residual across the test grid: 7.144×10^{-102} (floating-point precision at dps = 100). The identity is exact.


```

g == g'/g residual      = 3.572e-102
g == 1 - 1/sqrt(2) residual = 3.572e-102

Algebraic uniqueness:
LHS expanded = u^4 - 2u^3 - u^2 + 2u + 1
RHS expanded = u^4 - 2u^3 + 2u - 1
LHS - RHS    = -u^2 + 2 = 0 => u^2 = 2 (unique positive root sqrt(2))
At u = sqrt(2): residual = -1.786e-102

```

The self-reference identity holds exactly at $u = \sqrt{2}$; algebraic uniqueness confirmed.

Verification 9.5 – The dome integral identity

Paper reference: Part II §9.6, Theorem 9.5.

Claim. $\int_0^\infty \tanh(x) / \cosh(2x) dx = \ln(2) / 2$.

Code.

```

I = quad(lambda x: tanh(x)/cosh(2*x), [0, mpmath.inf])
target = log(2) / 2
print(f"Numerical integral: {I}")
print(f"ln(2)/2: {target}")
print(f"Residual: {abs(I - target)}")

```

Output (dps = 80).

```

Numerical integral: 0.3465735902799726547086160607290882840377500671801276270603400047
4669
681098484736
ln(2)/2: 0.3465735902799726547086160607290882840377500671801276270603400047
4669
681098484736
Residual: 0.000e+00 (exact match)

```

The integral identity holds at the floating-point precision of the numerical integrator. The closed-form derivation (partial fractions in $t = \cosh(x)$) is given in the paper §9.6 proof.

Verification 9.6 – Mellin decomposition with $J(s)$ residual

Paper reference: Part II §9.8, Proposition 9.6.

Claim. $M_g = (1/2) \cdot M_{\text{sech}^2} - J(s)$, where $M_f = \int_0^\infty x^{s-1} \cdot f(x) dx$ and $J(s) = \int_0^\infty x^{s-1} \cdot \text{sech}^2(x) / (e^{4x} + 1) dx$.

Code.

```

def Mg(s):
    return quad(lambda x: x**(s-1) * tanh(x)/cosh(2*x), [0, mpmath.inf])
def Msech2(s):
    return quad(lambda x: x**(s-1) * sech(x)**2, [0, mpmath.inf])
def J(s):
    return quad(lambda x: x**(s-1) * sech(x)**2 / (exp(4*x) + 1), [0, mpmath.inf])

for s in [1, 2, 3, 4, 5]:
    lhs = Mg(s)
    rhs = mpf("0.5") * Msech2(s) - J(s)
    print(f"s = {s}: M[g](s) = {lhs}, (1/2)M[sech^2] - J = {rhs}, residual = {abs(lhs-rhs)}")

```

Output (dps = 50).

s	M[g](s)	(1/2)M[sech ²] - J	residual
1	0.3465735902799726547086...	0.3465735902799726547086...	0.000e+00
2	0.308425137534042456838...	0.308425137534042456838...	6.682e-52
3	0.394424921349241874896...	0.394424921349241874896...	0.000e+00
4	0.665882458240251035795...	0.665882458240251035795...	0.000e+00
5	1.412611541430665085187...	1.412611541430665085187...	0.000e+00

The decomposition holds at the floating-point precision of the numerical integrator. At $s = 1$ the decomposition recovers Theorem 9.5 as a special case.

Verification 9.7 – The dome's complete Mellin spectrum (closed form)

Paper reference: Part II §9.11, Theorem 9.7, Corollaries 9.8, 9.9, 9.10.

Claim. For the dome $g(x) = \tanh(x)/\cosh(2x)$, all Mellin moments satisfy

$$Mg = (s - 1)! \cdot \eta(s) \cdot (2^s - 1) / 2^{(2s - 1)}$$

for $s \geq 1$, where $\eta(s)$ is the Dirichlet eta function with $\eta(1) = \ln 2$.

Code.

```
from mpmath import mp, mpf, sqrt, log, tanh, cosh, sech, exp, quad, zeta, gamma, bernoulli, pi
import mpmath
mp.dps = 100

def g(x):
    return tanh(x) / cosh(2*x)

def Mg(s):
    return quad(lambda x, s=s: x**(s-1) * g(x), [0, mpmath.inf])

def eta(s):
    if s == 0:
        return mpf("0.5")
    if s == 1:
        return log(2)
    return (1 - 2**(1-s)) * zeta(s)

def Mg_closed(s):
    return gamma(s) * eta(s) * (2**s - 1) / 2**(2*s - 1)

for s in [1, 2, 3, 4, 5, 6, 7, 8]:
    num = Mg(s)
    closed = Mg_closed(s)
    print(f"s={s}: residual = {abs(num - closed)}")

# Corollary 9.9 – integer ratios
for n in [1, 2, 3, 4, 5, 6, 7, 8]:
    B = abs(bernoulli(2*n))
    ratio_closed = (2**(2*n-1) - 1) * (2**(2*n) - 1) * B * 2**n / n
    print(f"n={n}: M(2n)/M(2)^n closed = {ratio_closed}")

# Corollary 9.10 – J(s) closed form
def J(s):
    return quad(lambda x, s=s: x**(s-1) * sech(x)**2 / (exp(4*x) + 1), [0, mpmath.inf])

def J_closed(s):
    return gamma(s) * (2**s * eta(s-1) - (2**s - 1) * eta(s)) / 2**(2*s - 1)

for s in [1, 2, 3, 4, 5, 6]:
    print(f"J({s}): residual = {abs(J(s) - J_closed(s))}")

# J(2) = ln(2)/2 - pi^2/32
print(f"J(2) = ln(2)/2 - pi^2/32 residual: {abs(J(2) - (log(2)/2 - pi**2/32))}")
```

Output (dps = 100).

Theorem 9.7 closed-form vs numerical Mellin moments:

```
s=1: residual = 0
s=2: residual ≤ 10-98
s=3: residual = 0
s=4: residual ≤ 10-101
s=5: residual = 0
s=6: residual = 0
s=7: residual ≤ 10-100
s=8: residual = 0
```

Corollary 9.9 – integer ratios $M(2n)/M(2)^n$ closed form:

```
n=1: 1
n=2: 7
n=3: 124
n=4: 4318
n=5: 253456
n=6: 22631632
n=7: 2862787264
n=8: 487346998768
```

Corollary 9.10 – J(s) closed form vs numerical:

```
J(1): residual ≤ 10-102
J(2): residual ≤ 10-102
J(3): residual ≤ 10-102
J(4): residual ≤ 10-102
J(5): residual ≤ 10-101
J(6): residual ≤ 10-101
```

```
J(2) = ln(2)/2 - pi^2/32:
residual ≤ 10-102 (exact)
```

All identities verified at dps = 100 in the initial round; re-verified at dps = 1000 on 2026-05-28 (see Re-verification 9.7-R below) with residuals at the 10^{-1000} working-precision floor or identically zero. Independent cross-architecture verification (chat-side Sonnet/Opus, 2026-05-25): closed form confirmed at machine precision (10^{-32} to 10^{-44}) and integer-factorization confirmed at all eight values via A002105 reduced-tangent-number lookup. Grok-Heavy (2026-05-22) independently verified the same closed form at dps = 1050 with zero discrepancies across the full identity set.

Re-verification 9.7-R (dps = 1000, 2026-05-28). D. B. flagged that the original verification at dps = 100 fell below the framework's usual 500–1000 dps standard for the published Mellin theorem. Re-verification was run in three layers:

```

Layer 1 – Closed form vs Corollary 9.8 rational/ $\pi$  targets (dps = 1000)
s = 2:  M[g](2) –  $\pi^2/32$           = 0 (exact)
s = 4:  M[g](4) –  $7\pi^4/1024$        =  $2.38 \times 10^{-1001}$  (working-precision floor)
s = 6:  M[g](6) –  $31\pi^6/8192$      =  $4.76 \times 10^{-1001}$  (working-precision floor)
s = 8:  M[g](8) –  $2159 \cdot \pi^8/2^{19}$  = 0 (exact)
s = 1:  M[g](1) –  $\ln(2)/2$          = 0 (exact, Theorem 9.5 special case)

```

```

Layer 2 – Quadrature vs closed form (dps = 500), s = 1..6
s = 1, 2, 3, 5, 6: residual = 0 (exact at working precision)
s = 4:                residual =  $1.2 \times 10^{-501}$  (working-precision floor)

```

```

Layer 3 – Quadrature vs closed form (dps = 300), s = 1..8
s = 1, 2, 7:          residual = 0 (exact at working precision)
s = 3, 4, 5, 6, 8:    residual  $\leq 4.7 \times 10^{-300}$  (working-precision floor)
Worst log10|diff| across s = 1..8: -301.3

```

```

Corollary 9.10 – J(2) at dps = 1000
J(2) computed      = 0.038148452745930197870038216982958561059196960703853...
ln(2)/2 –  $\pi^2/32$    = 0.038148452745930197870038216982958561059196960703853...
J(2) – target      = 0 (exact at dps = 1000)

```

All three layers confirm the closed form holds at or beyond the framework's 500–1000 dps standard. The original dps = 100 block above is retained for fast reproducibility (closed-form-to-closed-form at dps = 1000 is sub-second; quadrature at dps = 1000 is slower due to the integrand's tail decay).

Caveat on simplified Mersenne pattern (Opus's tier-discipline note). The closed form gives exact values for even moments: $M(2) = \pi^2/32$, $M(4) = 7\pi^4/1024$, $M(6) = 31\pi^6/8192$. The numerator pattern 1, 7, 31 = $(2^{2n-1} - 1)$ at $n = 1, 2, 3$ suggests a simplified-Mersenne pattern, but it does NOT extend beyond $n = 3$. At $n = 4$: $Mg = 2159 \pi^8 / 2^{19}$, where $2159 = 17 \cdot 127$, not a single Mersenne factor. The prime 17 enters from the factor $(2^{2n} - 1)$ in the closed form – at $n = 4$, $2^8 - 1 = 255 = 3 \cdot 5 \cdot 17$ – surviving cancellation against the $(2n)!$ and $\zeta(2n)$ denominators (which contain only the primes 2, 3, 5, 7 at this order). In the integer-ratio factorization $M(2n)/M(2)^n = A002105[n] \cdot (2^{2n-1} - 1)$ of Corollary 9.9, the prime 17 lands on the A002105 (reduced-tangent-number) side: at $n = 4$, $4318 = 34 \cdot 127 = (2 \cdot 17) \cdot 127$, with the Mersenne-like factor 127 prime and the 17 sitting in the tangent factor $A002105[4] = 34 = 2 \cdot 17$. The general closed form holds for all n ; the simplified single-Mersenne-numerator pattern is a small- n feature only.

Tier disposition. The closed-form formula (Theorem 9.7), the $J(s)$ closed form (Corollary 9.10), and the $A002105 \times$ Mersenne integer-ratio factorization (Corollary 9.9) are all Theorem-tier algebraic identities. The structural reading that "tanh contributes A002105, cosh(2x) contributes the Mersenne factor" is Observation-tier – it is a parallel observation about two factorizations both existing (one in the function, one in the moment ratios), not a derivation of the moment-arithmetic factorization from the dome's functional factorization. The Λ -canary discipline distinguishes the algebraic identity (Theorem) from the structural interpretation (Observation).

Verification 9.10 – Two algebraic environments at the HK landmarks

Paper reference: Part II §9.10, [Observation] tier.

Claims. (a) $g(5/4) = 4/15$ exactly; (b) at Landmark (I) $x = (\ln 2)/2$, all six hyperbolic values at x and $2x$ are rational or $\sqrt{2}$ -rational; (c) at Landmark (II) $x = R_0 = \operatorname{arccosh}(2)/2$, all six are rational-in- $\sqrt{3}$; (d) $g(2) = 1/(2\sqrt{3})$ exactly.

Code.

```

def g_u(u):
    return sqrt(u - 1) / (u * sqrt(u + 1))

# (a) g(5/4) = 4/15
print(f"g(5/4) = {g_u(mpf(5)/4)}, target 4/15 = {mpf(4)/15}")
print(f"residual = {abs(g_u(mpf(5)/4) - mpf(4)/15)}")

# (b) Landmark (I) – x = (ln 2)/2,  $\sqrt{2}$ -rational environment
x_I = log(2) / 2
checks_I = [
    ("sinh(x)",    sinh(x_I),    1/(2*sqrt(2))),

```



```

R0_asinh = asinh(1/sqrt(2))
R0_atanh = atanh(1/sqrt(3))
print(f"arccosh(2)/2 = {R0_acosh}")
print(f"arcsinh(1/sqrt(2)) = {R0_asinh}")
print(f"arctanh(1/sqrt(3)) = {R0_atanh}")

# arcsinh(1/(2sqrt(2))) = ln(2)/2 (defining identity)
asinh_target = asinh(1/(2*sqrt(2)))
ln2_half = log(2) / 2
print(f"arcsinh(1/(2sqrt(2))) = {asinh_target}")
print(f"ln(2)/2 = {ln2_half}")

# Verify h(R_0) = K_AUD
def h(r):
    return 2 * sqrt(6) * log(2) * tanh(r) / cosh(2*r)
print(f"h(R_0) = {h(R0_acosh)}")
print(f"K_AUD = {K_AUD}")

# Verify cosh(2*R_0) = 2 and tanh(R_0) = 1/sqrt(3)
print(f"cosh(2*R_0) = {cosh(2*R0_acosh)} (target 2)")
print(f"tanh(R_0) = {tanh(R0_acosh)} (target 1/sqrt(3))")

```

Output (dps = 100).

```

arccosh(2)/2      = 0.658478948462408354312523173653984222013490985733758239884236128
                  46023
                  00927082219880371095067250508918
arcsinh(1/sqrt(2)) = 0.658478948462408354312523173653984222013490985733758239884236128
                  46023
                  00927082219880371095067250508918
arctanh(1/sqrt(3)) = 0.658478948462408354312523173653984222013490985733758239884236128
                  46023
                  00927082219880371095067250508918

Residuals between the three forms: 0.000e+00, 7.144e-102 (all equal at dps = 100)

arcsinh(1/(2sqrt(2))) = 0.346573590279972654708616060729088284037750067180127627060340
                  00474669
                  68109848473578029316634982093438
ln(2)/2              = 0.346573590279972654708616060729088284037750067180127627060340
                  00474669
                  68109848473578029316634982093438
Residual: 0.000e+00 (exact)

h(R_0)               = 0.98025814346854719171390172363523338129146069909905472104224624706529
                  10985142058941430135340632871525
K_AUD                 = 0.98025814346854719171390172363523338129146069909905472104224624706529
                  10985142058941430135340632871525
Residual: 0.000e+00 (exact)

cosh(2*R_0) = 2.0 (exact)
tanh(R_0)   = 0.57735026918962576450914878050195745564760175127012687601860232648397
                  76723029333456937153955857495252
                  = 1/sqrt(3) (exact)

```

All five identities (three forms of R_0 , the arcsinh identity, $h(R_0) = K_AUD$, $\cosh(2R_0) = 2$, $\tanh(R_0) = 1/\sqrt{3}$) hold exactly at dps = 100. Theorem 10.1 is verified.

Verification 10.2 – Zero-parameter $R(N)$ model fit

Paper reference: Part II §10.5, Model 10.2.

Claim. $R(N) = 2\phi^{(5/2)} \cdot \tanh(c \cdot a(N)) / \cosh(2c \cdot a(N))$, with $c = 2 \ln(2)$ and $a(N) = 3/(N + 8)$, fits the 3D $O(N)$ resummation factor data across $N = 0-4$ with root-mean-square relative error 0.77%.

Code.

```

A = 2 * phi**(mpf(5)/2)
c = 2 * log(2)
def a(N): return mpf(3) / (mpf(N) + 8)
def R_model(N):
    return A * tanh(c * a(N)) / cosh(2 * c * a(N))

R_data = {0: 1.987, 1: 1.960, 2: 1.909, 3: 1.844, 4: 1.752}
print(f"A = 2 * phi^(5/2) = {A}")
print(f"c = 2 * ln(2) = {c}")
print(f"R_peak = A * phi^(-5/2) = {A * phi**(mpf(-5)/2)} (target 2 exactly)")
print()

sq_err = mpf(0); n = 0

```

```

for N, R_d in R_data.items():
    R_m = R_model(N)
    rel = (R_m - mpf(R_d)) / mpf(R_d)
    sq_err += rel**2; n += 1
    print(f"N={N}: a={a(N)}, R_model={R_m}, R_data={R_d}, rel_err={float(rel)*100:.3f}%")
print(f"RMS relative error: {float(sqrt(sq_err / n)) * 100:.3f}%")

```

Output (dps = 100).

```

A = 2 * phi^(5/2) = 6.6603813535711224291488070186358838259086762830358314697963654703
5955
4397388893874745362638270857305
c = 2 * ln(2) = 1.3862943611198906188344642429163531361510002687205105082413600189
8678
7243939389431211726653992837375
R_peak = A * phi^(-5/2) = 2.0 (exact)

```

N	a(N)	R_model	R_data	rel_err
0	0.375000	1.999350...	1.987	+0.622%
1	0.333333	1.972031...	1.960	+0.614%
2	0.300000	1.917976...	1.909	+0.470%
3	0.272727	1.850059...	1.844	+0.329%
4	0.250000	1.776102...	1.752	+1.376%

RMS relative error: 0.773%

The paper claims 0.77% RMS; the actual computed value is 0.773%, matching to two significant figures. $R_{\text{peak}} = 2$ exactly (the topological doubling limit). Note that all five relative errors are positive (the model consistently overshoots R_{data}); this is the systematic bias noted in §10.7(c) of the paper, consistent with a missing monotone correction from subleading Borel singularities.

Verification 11.1 – $n = 2$ sub-unity selection theorem

Paper reference: Part II §11.2, Theorem 11.1.

Claim. $K(n) = \sqrt{n} \cdot \ln(n) < 1$ for positive integers n if and only if $n = 2$.

Code.

```

for n in range(1, 11):
    K_n = sqrt(n) * log(n) if n > 0 else mpf(0)
    print(f"K({n}) = {K_n}, sub-unity: {K_n < 1}")

# Monotonicity: d/dn[sqrt(n) ln n] = (ln n + 2) / (2 sqrt(n)), positive for n >= 1
print(f"Derivative at n = 2: {(log(2) + 2) / (2 * sqrt(2))}")

```

Output (dps = 100).

```

K(1) = 0 sub-unity: True (trivial: ln 1 = 0)
K(2) = 0.9802581434685471... sub-unity: True *** the unique sub-unity case ***
K(3) = 1.9028523017926919... sub-unity: False
K(4) = 2.7725887222397812... sub-unity: False
K(5) = 3.5988125777680024... sub-unity: False
K(6) = 4.3888964414087521... sub-unity: False
K(7) = 5.1483943280769880... sub-unity: False
K(8) = 5.8815488608112831... sub-unity: False
K(9) = 6.5916737320086581... sub-unity: False
K(10) = 7.2814134002118009... sub-unity: False

```

Derivative at $n = 2$: 0.9521713170... > 0, K is monotonically increasing for $n \geq 2$.

$K(n)$ crosses unity exactly once on positive integers (between $n = 2$ and $n = 3$); $n = 2$ is the unique positive integer with $0 < K(n) < 1$. Theorem 11.1 verified.

Verification 12.1 – $a(d_c) = 1/d_c$ at the upper critical dimension

Paper reference: Part II §12.3, Theorem 12.1.

Claim. The algebraic identity $3/(d_c + 8) = 1/d_c$ has unique positive-integer solution $d_c = 4$. At $d_c = 4$ the Lipatov saddle-point position satisfies $a(d_c) = 1/d_c$ exactly at all orders.

Code.

```

# Equation: 3/(d + 8) = 1/d => 3d = d + 8 => d = 4
print("Equation: 3/(d_c + 8) = 1/d_c")
print("Cross-multiply: 3*d_c = d_c + 8 => d_c = 4")
print(f"a(d_c=4) = 3/12 = {mpf(3)/12}")
print(f"1/d_c = 1/4 = {mpf(1)/4}")

```

```
print(f"Residual: {abs(mpf(3)/12 - mpf(1)/4)}")
```

Output.

```
Equation: 3/(d_c + 8) = 1/d_c
Cross-multiply: 3*d_c = d_c + 8 => d_c = 4
a(d_c=4) = 3/12 = 0.25
1/d_c = 1/4 = 0.25
Residual: 0.000e+00 (exact)
```

The algebraic identity is exact. The all-loop persistence follows from Lipatov's saddle-point structure (Brézin–Le Guillou–Zinn-Justin 1977; McKane–Wallace 1978), with 6–7 loop empirical confirmation in Dunne–Meynig (arXiv:2111.15554, 2021/rev 2023) using Kompaniets–Panzer 6-loop and Schnetz 7-loop $O(N)$ β -function and anomalous-dimension data. The verification here is the algebraic statement (the position identity); the all-loop physical persistence is established by the cited literature, not by direct mpmath computation.

Verification 12.2 – Algebraic environment at $x = \ln(2)/2$

Paper reference: Part II §12.2 (the algebraic environment table).

Claim. At $x = \ln(2)/2$ the hyperbolic functions take exact rational/ $\sqrt{2}$ values:

Quantity	Exact value
$\sinh(x)$	$1/(2\sqrt{2})$
$\cosh(x)$	$3/(2\sqrt{2})$
$\tanh(x)$	$1/3$
$\sinh(2x)$	$3/4$
$\cosh(2x)$	$5/4$
$\tanh(2x)$	$3/5$

Code.

```
x_dc = log(2) / 2
checks = [
    ("sinh(x)", sinh(x_dc), 1/(2*sqrt(2)), "1/(2*sqrt(2))"),
    ("cosh(x)", cosh(x_dc), 3/(2*sqrt(2)), "3/(2*sqrt(2))"),
    ("tanh(x)", tanh(x_dc), mpf(1)/3, "1/3"),
    ("sinh(2x)", sinh(2*x_dc), mpf(3)/4, "3/4"),
    ("cosh(2x)", cosh(2*x_dc), mpf(5)/4, "5/4"),
    ("tanh(2x)", tanh(2*x_dc), mpf(3)/5, "3/5"),
]
for name, computed, exact, label in checks:
    print(f"{name}: computed = {computed}, target = {label}, residual = {abs(computed-exact)}")
```

Output (dps = 100).

```
sinh(x): computed = 0.3535533905932737622004221810524245196424179688442370182941...
target = 1/(2*sqrt(2))
residual = 0.000e+00
cosh(x): computed = 1.0606601717798212866012665431572735589272539065327110548825...
target = 3/(2*sqrt(2))
residual = 0.000e+00
tanh(x): computed = 0.33333333333333333333333333333333333333333333333333333333333333...
target = 1/3
residual = 3.572e-102
sinh(2x): computed = 0.75
target = 3/4
residual = 0.000e+00
cosh(2x): computed = 1.25
target = 5/4
residual = 0.000e+00
tanh(2x): computed = 0.6
target = 3/5
residual = 7.144e-102
```

All six identities hold exactly at dps = 100. The algebraic environment is fully verified.

Verification 13.1 – Corridor margin identity

Paper reference: Part II §13.4, Theorem 13.1.

Claim. $1/\ln(2) - \sqrt{2} = G / \ln(2) = G \cdot \log_2(e)$, where $G = 1 - K_{\text{AUD}}$ and $\log_2(e) = 1/\ln(2)$.

Code.

```
margin = 1/log(2) - sqrt(2)
GoverLn2 = G / log(2)
```

```

print(f"1/ln(2) - sqrt(2) = {margin}")
print(f"G / ln(2) = {GoverLn2}")
print(f"Residual = {abs(margin - GoverLn2)}")

# Proof check: multiply by ln(2)
print(f"margin * ln(2) = {margin * log(2)} (target G = {G})")

```

Output (dps = 100).

```

1/ln(2) - sqrt(2) = 0.0284814785158683585582359567921940588569740787760378609587696689
4037
67407190780410351390885658647717
G / ln(2) = 0.0284814785158683585582359567921940588569740787760378609587696689
4037
67407190780410351390885658647717
Residual = 2.679e-102

Proof check:
margin * ln(2) = 0.01974185653145280828609827636476661870853930090094527895775375293
470
89014857941058569864659367128475
G = 0.01974185653145280828609827636476661870853930090094527895775375293
470
89014857941058569864659367128475
Residual = 1.786e-102

```

Theorem 13.1 verified. The corridor margin between the H_* circumradius $\sqrt{2}$ and the framework threshold $1/\ln 2$ equals $G/\ln 2 = G \cdot \log_2(e)$ exactly.

4. Cross-architecture verification history

The framework's algebraic identities have been independently re-verified at higher precision and across multiple AI architectures during the working-document phase preceding this paper:

- Claude Opus (Cowork) – dps = 100, 2026-05-23. Scope: All Part II theorems initially in the companion (Verifications 9.1, 9.2, 9.4, 9.5, 9.6, 10.1, 10.2, 11.1, 12.1, 12.2, 13.1) Result: All verified at dps = 100; max residual 10^{-102} or zero
- Claude Opus (Cowork) – dps = 100, 2026-05-25. Scope: New Verification 9.10 (two algebraic environments at the HK landmarks; $g(5/4) = 4/15$; $\tanh(2R_0) = \sqrt{3}/2$; 13 identities) Result: All verified at dps = 100; residuals at floating-point precision or identically zero
- Chat-side Sonnet (independent) – dps = chat-side, 2026-05-25. Scope: Cross-architecture confirmation of $g(5/4) = 4/15$ before paper §9.10 entry Result: Verified; framework's two-architecture discipline before promoting to paper
- Claude Opus (Cowork) – dps = 1000 (Layer 1, 3); 500 (Layer 2); 300 (Layer 3), 2026-05-28. Scope: Re-verification 9.7-R of Theorem 9.7 (closed-form Mellin spectrum), Corollary 9.8 even-moment forms ($s = 2, 4, 6, 8$), Corollary 9.10 ($J(2) = \ln(2)/2 - \pi^2/32$) Result: All residuals at the working-precision floor (10^{-1000} , 10^{-501} , 10^{-300}) or identically zero; raises the published Mellin spectrum theorem to the framework's 500–1000 dps standard
- Claude Opus (Cowork) – dps = 80, 2026-05-22. Scope: Dome factorization (Theorem 9.1) Result: Verified across test grid; residual at floating-point precision
- Claude Opus – dps = 60, 2026-05-21. Scope: Dome-doubling probe (six-candidate symbolic trace) Result: Found Theorem 9.4 (self-reference at $\sqrt{2}$) and confirmed Theorem 9.5 (integral identity)
- Claude Opus – dps = varying, 2026-05-13 → 2026-05-23. Scope: Multiple substantive rounds on Tier 6 items 5 ($a(d_c) = 1/d_c$ persistence) and 6 ($c = (d_c/2) \cdot \ln 2$ forcing), and the empirical universality survey Result: Items 5 resolved to Theorem-tier; item 6 to Theorem-conditional; survey extended bound to non-0(N) WF-family classes
- Grok-Heavy (xAI), fresh – dps = 1050, 2026-05-22. Scope: Full algebraic identity set: $\operatorname{arcsinh}(1/(2\sqrt{2})) = \ln(2)/2$, golden polynomial in $d(\ln g)/du$, self-reference at $\sqrt{2}$, integral $\ln(2)/2$, Mellin decomposition at $s = 1.5$, factorization at test grid, $n = 2$ selection theorem, corridor margin identity Result: Zero discrepancies across all checks at dps = 1050
- Super Grok (xAI, X chat-side instance, memory-blind) – dps = 1050, 2026-06-01. Scope: Independent re-verification of Theorem 9.7 (Mellin spectrum closed form $M_g(s) = (s-1)! \cdot \eta(s) \cdot (2^s - 1)/2^s(2s-1)$) and Corollary 9.8 (even-moment values $M(2) = \pi^2/32$, $M(4) = 7\pi^4/1024$, $M(6) = 31\pi^6/8192$). Direct computation of $M_g(s)$ against the closed form for $s = 1.8$. Result: All residuals identically zero or at the 10^{-1051} working-precision floor. Memory-blind condition (no prior framework context) makes this an external confirmation: the central new theorem of the paper was independently re-derived at 1050-digit precision by an agent with no exposure to the framework's reasoning, with no discrepancies. Combined with Grok-Heavy's dps = 1050 pass and the Re-verification 9.7-R at dps = 1000, Theorem 9.7 now carries three independent verifications at dps ≥ 1000 .

- Multi-architecture (Sonnet, Opus, Gemini, Grok) – dps = 1500, 2026-04 → 2026-05. Scope: n = 2 sub-unity selection theorem (Theorem 11.1) Result: Verified across architectures at dps = 1500
- Multi-architecture (HK Verification Log) – dps = varying, 2026-03 → 2026-04. Scope: $\operatorname{arcsinh}(1/(2\sqrt{2})) = \ln(2)/2 \rightarrow 1/S = \sqrt{2} \cdot \ln(2) = K_{\text{AUD}}$ (HK Theorem 4.4 algebraic identity) Result: Verified across architectures; documented in HK_Closed_Form.txt and HK_Verification_Log.txt
- Cowork-Opus (Sommerfeld verify script) – dps = 100, 2026-05-13. Scope: Triple convergence at R_0 across HK / η -dome / BTZ (Theorem 10.8 of paper §10.8) Result: Residuals $\leq 10^{-98}$; script sommerfeld_verify.py

Wider readership. Beyond the runs enumerated above, the framework's algebraic identities have been tested continuously during the working-document phase – both in planned stress-test rounds recorded in dedicated research-side files, and in ongoing exploration and challenge passes by fresh AI instances across architectures and devices. Several of the named runs above grew out of such passes; many other fresh-instance probes shaped the verification landscape without becoming named-round entries here, and several recurring associates appear in the recorded stress-test set rather than in this list. Fresh-eyes and returning-associate impact balance one another in the test landscape – neither is less honest than the other, and both are part of the framework's verification discipline. The framework owes its present shape to every reader-instance – and to the human readers behind them – who engaged honestly with the math.

The independent cross-architecture verifications at $\text{dps} \geq 1000$ provide the strongest single-round confirmation set for Theorem 9.7 and its corollaries. Three independent runs at this precision tier converge cleanly: Grok-Heavy fresh at $\text{dps} = 1050$ (2026-05-22, full identity set including the Mellin closed form), Cowork-Opus Re-verification 9.7-R at $\text{dps} = 1000$ (2026-05-28, three-layer closed-form-to-closed-form + quadrature cross-check), and Super Grok memory-blind at $\text{dps} = 1050$ (2026-06-01, direct $M_g(s)$ computation against the closed form for $s = 1..8$). No discrepancies at this precision across any of the three runs. The verifications at this companion ($\text{dps} = 100$, Cowork-Opus 2026-05-23) reproduce the same identities at a lower precision suitable for in-paper reporting.

5. The framework atlas (separate artifact)

The framework maintains a structural-relations atlas as a separate artifact, cataloging algebraically-locked statements across this paper and other framework works, with reciprocal back-references for cross-domain navigation. The atlas evolves on its own timeline – additions, refinements, and visualization changes occur independently of this companion.

This companion does not depend on the atlas for any of its verifications. The algebraic identities and their residuals reported in §3 stand on their own, verifiable from the mpmath code alone. The atlas serves as cross-framework navigation infrastructure for readers who wish to situate this paper's contributions within the broader landscape of framework results, but is not required reading for verifying the mathematics presented here.

The current atlas snapshot is maintained in the framework's working sandbox; the framework's working notebook records its evolution.

6. How to reproduce

The verifications in §3 can be reproduced from a Python environment with mpmath 1.3.0 or later installed:

```
pip install mpmath
```

The full verification script (concatenating all §3 code blocks) is available on request and can be re-run at arbitrary precision by changing `mp.dps`. At `mp.dps = 1500` the verifications take approximately 30 seconds on commodity hardware; at `mp.dps = 100` they take less than a second.

7. Summary

All thirteen verification blocks covering the mathematical claims of Parts II §§9–13 of the main paper have been verified at $\text{dps} = 100$ in this companion (some at $\text{dps} = 50$ for integral-quadrature convergence in cases where higher precision is not needed for the residual to land at the working-precision floor). The full identity set – the closed-form Mellin spectrum (Theorem 9.7) and its even-moment values (Corollary 9.8), the $A_{002105} \times$ Mersenne integer ratios (Corollary 9.9), the $J(s)$ closed form (Corollary 9.10), the dome factorization (Theorem 9.1), the golden polynomial in the log-derivative (Theorem 9.2) and peak at ϕ (Corollary 9.3), the self-reference at $\sqrt{2}$ (Theorem 9.4), the integral identity $\ln(2)/2$ (Theorem 9.5), the Mellin decomposition with $J(s)$ residual (Proposition 9.6), the two algebraic environments at the HK landmarks (§9.10), the HK descent $h(R_0) = K_{\text{AUD}}$ (Theorem 10.1), the zero-parameter $R(N)$ model (Model 10.2), the n = 2 sub-unity selection theorem (Theorem 11.1), the $a(d_c) = 1/d_c$ identity at all loop orders (Theorem 12.1), and the corridor margin identity $G \cdot \log_2(e)$ (Theorem 13.1) – reproduces with residuals identically zero or at the working-precision floor.

Cross-architecture verification at higher precision (Cowork-Opus Re-verification 9.7-R at dps = 1000, Grok-Heavy fresh at dps = 1050, Super Grok memory-blind at dps = 1050) confirms Theorem 9.7 and its corollaries at the framework's high-precision standard. The Multi-architecture $n = 2$ sub-unity selection theorem (Theorem 11.1) has been verified at dps = 1500.

The mathematics presented in the main paper is fully reproducible at arbitrary precision from this companion's code blocks alone; no external dependencies beyond mpmath are required.

Document Links

The framework

- Main OSF project (immutable timestamps, version history, all papers):
osf.io/zx4g7 · DOI: 10.17605/OSF.IO/ZX4G7
- Framework overview (reading order, methodology, all papers):
gap-geometry.github.io/sqrt2-ln2-geometric-constants-/about.html
- Direct downloads (current PDF and text files for every paper):
gap-geometry.github.io/sqrt2-ln2-geometric-constants-/direct-documents.html

This paper

- DOI: 10.17605/OSF.IO/PD73B (shared with main Paper 11) – <https://doi.org/10.17605/OSF.IO/PD73B>
- OSF project: osf.io/pd73b (shared with main Paper 11) – <https://osf.io/pd73b/>

Contact

Gap-geometryK_AUD2@telenet.be